LOUKIK NAIK

loukiknaik@gmail.com | linkedin/loukiknaik | portfolio.loukik.dev | +18582419592 | **San Francisco, CA**

Education

University of California, San Diego — California, USA
Master of Science (MS) in Computer Science (GPA: 3.97/4.0) — *December 2024*

University of Mumbai — Mumbai, India
Bachelor of Engineering in Computer Engineering (GPA: 9.71/10.0) — *May 2023*

Work Experience

Software Engineer — Palo Alto, CA
Eudia (Legal AI Startup, Series A) — *September 2025 – Present*

- Owned the **Mergers and Acquisitions** backend service end-to-end, building and operating the core infrastructure for an **agent-based AI system** for large-scale legal document analysis in production.
- Partnered closely with **lawyers** and **product managers** to design and ship product features, translating legal workflows and requirements into scalable platform capabilities.
- Designed and shipped backend APIs that act as the control plane for **AI agent execution**, managing document ingestion, run orchestration, result persistence, and retrieval of structured, tabular outputs.
- Re-architected agent orchestration by migrating from **Airflow** to **Temporal**, enabling reliable execution of long-running, stateful agents with retries, idempotency, and deterministic behavior.
- Built backend capabilities for multiple **AI agents** including flexible document analysis, configurable PII redaction, Contract Comparison and tabular review with **schema-based inference** and **cell-level operations** backed by structured logging, workflow tracing, telemetry, and **Kubernetes**-native monitoring.
- Built workflows coordinating multiple **services**, **queues**, and **workers**, addressing production challenges such as duplicate execution, inconsistent state across retries, and downstream dependency failures.
- Designed and managed **Postgres** database schemas and migrations using **Alembic**, delivering backward-compatible schema changes to support new agent capabilities.
- Converted the service to a **multi-tenant architecture**, enforcing **tenant-aware routing**, strict isolation, and safe rollouts across live customers.
- Managed the rollout of new backend features across **staging, internal, and production** environments using **Kubernetes, Helm, and ArgoCD**, supporting multiple customer deployments.

Machine Learning Engineer — San Francisco, USA
Plainsight Technologies (B2B Computer Vision Startup, Seed Round) — *January 2025 - August 2025*

- Built and deployed real-time image segmentation APIs for SAM and MobileSAM using **FastAPI**, with Dockerized model serving and automated infrastructure and deployments via **Terraform** and CI/CD.
- Created a **cron-triggered retraining pipeline** that detects newly labeled data and launches **Vertex AI** training jobs, enabling continuous model updates.
- Designed an **OCR model evaluation framework** utilizing **edit distance, confidence scoring**, and **text similarity** metrics to automate quality checks pre-deployment.
- Designed and implemented a robust **integration testing framework** to validate end-to-end performance of the model generation pipelines for **segmentation, object** and **keypoint detection** tasks.

MLE Intern — San Diego, USA
Plainsight Technologies — *June 2024 - December 2024*

- Automated video inference workflows using **Google Cloud Functions** to trigger ML pipelines on GCS uploads.
- Deployed these pipelines on **Google Kubernetes Engine (GKE)** with CI/CD for scalable, reliable processing.

Project

Surfstore: A Distributed File Storage System | *Software Development, Computer Networks*

- Built a **horizontally scalable** file storage system in **Go**, storing data across multiple block and metadata servers.
- Built a **gRPC-based client–server architecture** for efficient communication between storage nodes.
- Implemented the **Raft consensus algorithm** to ensure consistency and fault tolerance across metadata servers.

Technical Skills

Software Engineering & Systems: Python, C++ (STL), JavaScript, SQL, Bash, HTML/CSS; Data Structures, OOP, System Design, Distributed Systems, Cloud Computing, Design Patterns; FastAPI, Git

Machine Learning & Infrastructure: LLMs, NLP; PyTorch, TensorFlow, Pandas, Scikit-Learn, LangGraph, LangChain, TTS; Kubernetes, Docker, Kubeflow, Temporal, CI/CD Pipelines, Grafana, Prometheus, Signoz, ArgoCD